

CS 188: Artificial Intelligence Notes (for Reinforcement Learning)

Pranay Pasula

July 21, 2019

Intro

The following questions test a basic understanding of Markov decision processes and reinforcement learning in the context of CS 188: Artificial Intelligence at UC Berkeley.

Sorry for any typos.

Lecture 8: Markov Decision Processes I

1. How is an MDP defined?
2. How does an MDP relate to a general search problem?
3. What is Markovian about MDPs?
4. What does this imply about how one should design an MDP?
5. What is usually the main goal in a deterministic search problem?
6. What is usually the main goal in an MDP?
7. What is a policy?
8. What does an explicit policy define? Explain what this is.
9. How does expectimax differ from finding an optimal policy?
10. How does expectimax relate to finding an optimal policy?
11. What is an MDP search tree?
12. How does an MDP search tree relate to an expectimax search tree?
13. How does an MDP search tree differ from an expectimax search tree?
14. What is discounting?
15. What issue may arise without discounting? What other condition must hold for this to be possible?
16. What are stationary preferences?
17. If we assume stationary preferences, what is implied about the ways to define utilities?

18. What is a potential issue if an MDP can last forever (and accrue infinite rewards)? What are three ways to address this? What are the two most common ways to address this? How does the third way relate to the probability of the MDP terminating?
19. What are the three optimal values of interest when solving MDPs? Explain the meaning of each one.
20. How many Q-values are associated with each state?
21. How do Q-values relate to values?
22. How can we find the value of a state?
23. What is the formula for the value of a state?
24. What is the formula for the Q-value of a pair (s, a) ?
25. What are two potential issues with using just expectimax to solve an MDP (hint: consider an infinite-length MDP)? How can we address these?
26. What are time-limited values? How can we quickly compute $V_i(s)$ for any $i \in [0, 1, \dots, k]$? In general, for what value of i is $V_i(s) = V^*(s)$?
27. Explain the value iteration algorithm.
28. What is its computational complexity?
29. Is the solution unique?
30. Is the solution optimal?
31. What is the main problem with value iteration?
32. Does value iteration always converge?
33. How does the convergence of the values relate to the convergence of the policy?
34. Does value iteration converge to the optimal policy?
35. How does adjusting γ affect the convergence rate?

Lecture 9 - Markov Decision Processes II

1. What is policy evaluation?
2. What is a fixed policy?
3. Explain two ways to find the values for a fixed policy.
4. How does this compare to value iteration?
5. What is the complexity of the iterative approach to policy evaluation?
6. For the non-iterative approach, what kind of matrix is usually involved? (Hint: usually enables a computer to find the solution quickly)
7. Given the optimal value for each state, can we determine the optimal policy without any calculations?
8. Explain how to determine the optimal policy given optimal value for each state.
9. Explain how to determine the optimal policy given the Q-values for each state.
10. State three problems with value iteration.
11. How does the rate of changes of "max" actions differ for early, middle, and late stages of value iteration?
12. What are the two main steps of policy iteration? Explain each.

13. Does policy iteration converge to the optimal policy? Does it converge to the optimal values? Explain why.
14. Why use policy iteration instead of value iteration? For what type of MDP is policy iteration especially preferred over value iteration?
15. Does policy iteration result in a better solution than value iteration?
16. Explain how to find the optimal policy using value iteration.
17. Explain how to find the optimal values using policy iteration.
18. When does policy iteration terminate? How does $\gamma \in (0, 1)$ affect this?
19. What is a multi-armed bandit?
20. What kind of rewards do choices in the multi-armed bandit problem usually have? How does this affect the corresponding MDP?
21. Explain how to find the optimal policy of a multi-armed bandit MDP given all of the details of the MDP. What is this called? State why this is advantageous to having only partial information of the MDP.
22. Explain how to find the optimal policy of a multi-armed bandit MDP given only partial details of the MDP. What is this called?

Lecture 10 - Reinforcement Learning I

1. Describe the basic feedback loop of reinforcement learning.
2. What feedback is received by the agent?
3. What is the goal of the agent?
4. What does the agent learn from?
5. What is the goal of reinforcement learning?
6. How does a reinforcement learning problem differ from an MDP problem?
7. What is model-based learning? Explain how to do it.
8. How do the estimated values in model-based learning act as the number of experiences approaches infinity?
9. What is a disadvantage of model-based learning?
10. What is model-free learning?
11. What is an advantage of model-free learning over model-based learning?
12. What is passive reinforcement learning? What is the goal? How does this differ from offline planning?
13. What is direct evaluation? What are its advantages and disadvantages?
14. What is policy evaluation? Can we use it to solve the passive RL problem? Why or why not?
15. What is sample-based policy evaluation? Why may it not be suitable to solve the passive RL problem?
16. What is temporal difference learning? When do updates occur? How do updates occur? What info does this algorithm take into account that direct evaluation neglects?

17. What type of average does temporal difference learning use? How does this average behave? How does the parameter of this average affect the behavior? How can we adjust this type of average to promote convergence?
18. How does temporal difference learning relate to Bellman updates?
19. What is the main problem with temporal difference learning? What should we do to avoid this?
20. What is active reinforcement learning? What is its goal?
21. How does active RL differ from passive RL?
22. What is the fundamental trade-off in active reinforcement learning?
23. Why can't we use value iteration to solve the active RL problem?
24. What is Q-value iteration? Why can't we use this to solve the active RL problem?
25. What is Q-learning? How does it relate to temporal difference learning? What conditions are required for it to converge to the optimal policy?
26. What is off-policy learning? Is Q-learning considered on-policy or off-policy? Explain why.

Lecture 11 - Reinforcement Learning II

1. How does our representation of an RL problem compare to an MDP? What the goal in RL?
2. Explain the methods used to solve for $V^*(s)$, $Q^*(s)$, and $\pi^*(s)$ and to evaluate a fixed policy for an MDP. What type of learning is this?
3. What are the two classes of methods used to solve for $V^*(s)$, $Q^*(s)$, and $\pi^*(s)$ and to evaluate a fixed policy in an RL problem? What are the specific methods used to do both of these tasks? How do the two classes of methods compare to one another? What type of learning are these?
4. What is Q-value iteration? What is Q-learning?
5. What is the exploration versus exploitation trade-off?
6. What is the simplest way to choose actions? What are the issues with this method?
7. What is an exploration function? State a basic example of one. How does this affect the Q-value update?
8. How does the bonus of a particular state affect the Q-values of other states?
9. How is regret defined? Why isn't learning the optimal policy enough to minimize regret? How does the regret of random exploration (ϵ -greedy) compare to the regret of exploration functions? Do these methods converge to a particular policy? If so, do they converge to the optimal policy?
10. What is an issue with finding the exact Q-value for every state-action pair (I.e., $Q(s, a)$ for every (s, a) possible)? Explain feature-based learning and how it addresses this issue. What is a q function? What type of function is this? Name an advantage and a disadvantage of representing $Q(s, a)$ with a q function. What class of Q-learning is feature-based learning? How is the learning process affected by using this class of Q-learning instead of finding exact Q-values for all of the state-action pairs.

11. Explain approximate Q-learning. What is the mathematical justification for the update step involved?
12. What is policy search? Why use it? In this context, what is the goal of Q-learning? What is the goal of action selection? What trade-off is involved between these two? What is the goal of policy search?